


```
import pandas as pd
df = pd.read_excel('CTG.xls','Data2')
df.head()
```



	b	e	AC	FM	UC	DL	DS	DP	DR	Unnamed: 9	...	E	AD	DE	LD	FS	SUSP	Unnamed: 42	CLASS	Unnamed: 44	NSP
0	240.0	357.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	NaN	...	-1.0	-1.0	-1.0	-1.0	1.0	-1.0	NaN	9.0	NaN	2.0
1	5.0	632.0	4.0	0.0	4.0	2.0	0.0	0.0	0.0	NaN	...	-1.0	1.0	-1.0	-1.0	-1.0	-1.0	NaN	6.0	NaN	1.0
2	177.0	779.0	2.0	0.0	5.0	2.0	0.0	0.0	0.0	NaN	...	-1.0	1.0	-1.0	-1.0	-1.0	-1.0	NaN	6.0	NaN	1.0
3	411.0	1192.0	2.0	0.0	6.0	2.0	0.0	0.0	0.0	NaN	...	-1.0	1.0	-1.0	-1.0	-1.0	-1.0	NaN	6.0	NaN	1.0
4	533.0	1147.0	4.0	0.0	5.0	0.0	0.0	0.0	0.0	NaN	...	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	NaN	2.0	NaN	1.0

5 rows × 46 columns

```
df = df.drop(df.columns[df.columns.str.contains('Unnamed', na=False)], axis=1)
df = df.drop(df.columns[0:9], axis=1)
```

```
df.head()
```

	LB	AC.1	FM.1	UC.1	DL.1	DS.1	DP.1	ASTV	MSTV	ALTV	...	C	D	E	AD	DE	LD	FS	SUSP	CLASS	NSP
0	120.0	0.000000	0.0	0.000000	0.000000	0.0	0.0	73.0	0.5	43.0	...	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	1.0	-1.0	9.0	2.0
1	132.0	0.006380	0.0	0.006380	0.003190	0.0	0.0	17.0	2.1	0.0	...	-1.0	-1.0	-1.0	1.0	-1.0	-1.0	-1.0	-1.0	6.0	1.0
2	133.0	0.003322	0.0	0.008306	0.003322	0.0	0.0	16.0	2.1	0.0	...	-1.0	-1.0	-1.0	1.0	-1.0	-1.0	-1.0	-1.0	6.0	1.0
3	134.0	0.002561	0.0	0.007682	0.002561	0.0	0.0	16.0	2.4	0.0	...	-1.0	-1.0	-1.0	1.0	-1.0	-1.0	-1.0	-1.0	6.0	1.0
4	132.0	0.006515	0.0	0.008143	0.000000	0.0	0.0	16.0	2.4	0.0	...	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	-1.0	2.0	1.0

5 rows × 33 columns

```
df.isnull().sum()
```

LB	2
AC.1	2
FM.1	2
UC.1	2
DL.1	1
DS.1	1
DP.1	1
ASTV	2
MSTV	2
ALTV	2
MLTV	2
Width	2
Min	2
Max	2
Nmax	2
Nzeros	2
Mode	2
Mean	2
Median	2
Variance	2
Tendency	2
A	1
B	1
C	1
D	1
E	1
AD	1
DE	1
LD	1
FS	1
SUSP	1
CLASS	2
NSP	2
dtype:	int64

```
df=df.dropna()
df.isnull().sum()
```

LB	0
AC.1	0
FM.1	0
UC.1	0
DL.1	0
DS.1	0
DP.1	0
ASTV	0
MSTV	0
ALTV	0
MLTV	0

```
Width      0
Min        0
Max        0
Nmax       0
Nzeros     0
Mode       0
Mean       0
Median     0
Variance   0
Tendency   0
A          0
B          0
C          0
D          0
E          0
AD         0
DE         0
LD         0
FS         0
SUSP       0
CLASS      0
NSP        0
dtype: int64
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
correlation_with_target = df.corr()[['NSP']]
print("Correlation with Target Column:")
print(correlation_with_target)
```

```
sns.set(style="whitegrid")
plt.figure(figsize=(10, 6))
sns.barplot(x=correlation_with_target.index, y=correlation_with_target['NSP'])
plt.xticks(rotation=90)
plt.title(f'Correlation with NSP')
plt.show()
```

```
Correlation with Target Column:
      NSP
LB      0.148151
AC.1    -0.363849
FM.1     0.087933
UC.1    -0.203824
DL.1     0.062702
DS.1     0.135629
DP.1     0.488277
ASTV     0.471191
MSTV    -0.103382
ALTV     0.426146
MLTV    -0.226797
Width   -0.068789
Min      0.063175
Max     -0.045265
Nmax    -0.023666
Nzeros  -0.016682
.....

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import chi2_contingency

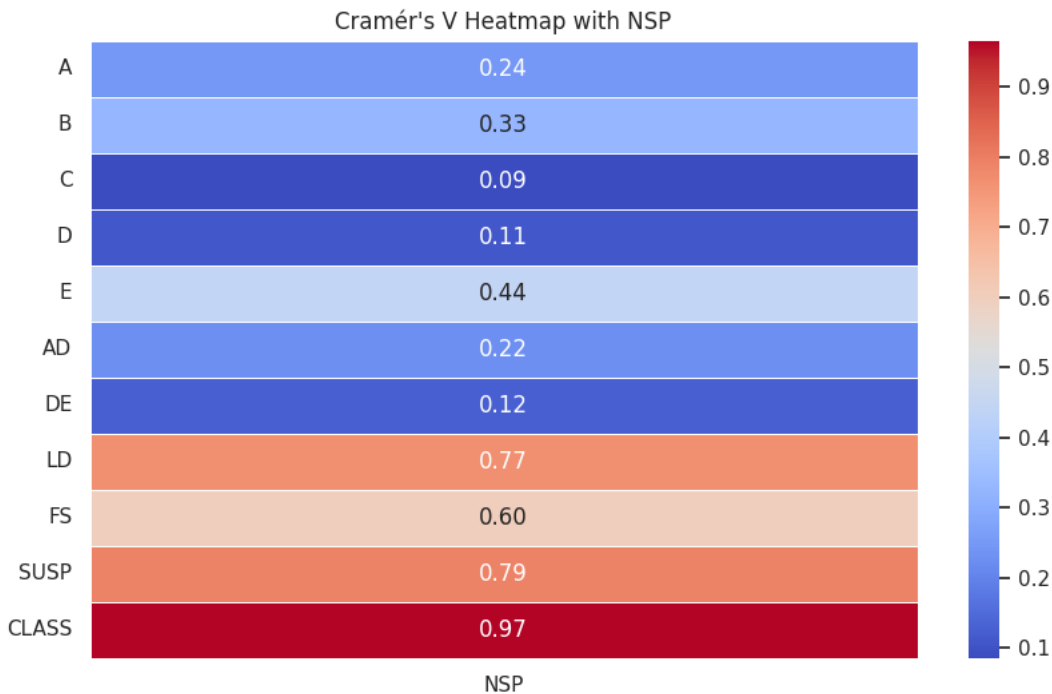
categorical_features = ['A','B','C','D','E', 'AD','DE','LD','FS','SUSP','CLASS']
target_variable = 'NSP'

contingency_tables = pd.DataFrame(index=categorical_features, columns=[target_variable])

for feature in categorical_features:
    contingency_table = pd.crosstab(df[feature], df[target_variable])
    chi2, _, _, _ = chi2_contingency(contingency_table)
    n = np.sum(contingency_table.sum())
    cramers_v = np.sqrt(chi2 / (n * (min(contingency_table.shape) - 1)))
    contingency_tables.at[feature, target_variable] = cramers_v


contingency_tables = contingency_tables.apply(pd.to_numeric)

plt.figure(figsize=(10, 6))
sns.heatmap(contingency_tables, annot=True, cmap="coolwarm", fmt=".2f", linewidths=.5)
plt.title(f'Cramér's V Heatmap with {target_variable}')
plt.show()
```




```
plt.figure(figsize=(16, 8))
cols = [ 'DP.1', 'ASTV','ALTV']
sns.boxplot(data=df[cols], orient="h", palette="Set2")
plt.title('Box Plots for Outliers in Columns')
plt.show()
```

Box Plots for Outliers in Columns



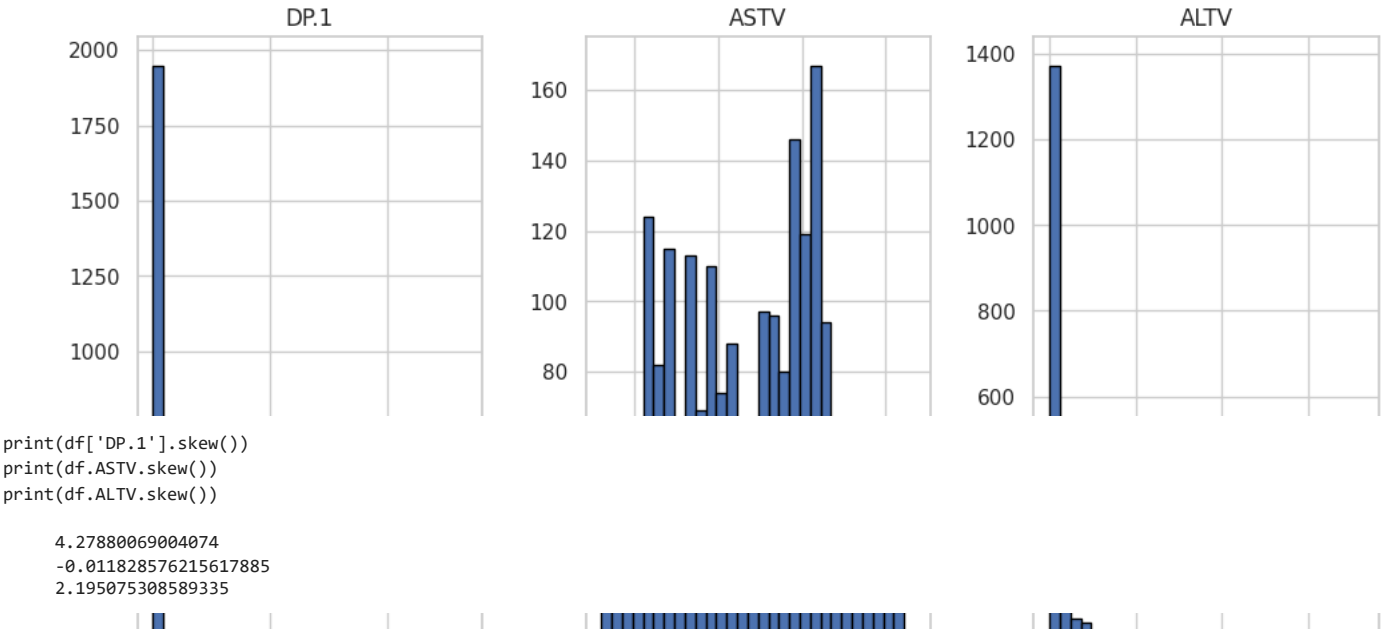
```
sns.set(style="whitegrid")
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))
sns.countplot(x='LD', data=df, ax=axes[0], palette='Set2')
sns.countplot(x='FS', data=df, ax=axes[1], palette='Set2')
sns.countplot(x='CLASS', data=df, ax=axes[2], palette='Set2')
plt.suptitle('Count Plots of Categorical Features', y=1.02)
plt.tight_layout()
plt.show()
```

Count Plots of Categorical Features



```
sns.set(style="whitegrid")
df[['DP.1', 'ASTV', 'ALTV']].hist(bins=30, figsize=(12, 6), edgecolor='black', layout=(1, 3))
plt.suptitle('Histograms of Continuous Features', y=1.02)
plt.show()
```

Histograms of Continuous Features




```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder
X = df[['LD', 'FS', 'CLASS',
        'DP.1', 'ASTV', 'ALTV']]
y = df['NSP']

X = pd.get_dummies(X, columns=['LD', 'FS', 'CLASS'], drop_first=True)
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
clf_rf = RandomForestClassifier(n_estimators=100, random_state=42)
clf_rf.fit(X_train, y_train)
y_pred = clf_rf.predict(X_test)
accuracy_rf = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy_rf:.2f}')
print('\nClassification Report:\n', classification_report(y_test, y_pred))
```

Accuracy: 0.98

Classification Report:				
	precision	recall	f1-score	support
0	0.99	0.99	0.99	333
1	0.95	0.94	0.94	64
2	1.00	0.97	0.98	29
accuracy			0.98	426
macro avg	0.98	0.97	0.97	426
weighted avg	0.98	0.98	0.98	426

```
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cf= confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix for Random Forest')
```



```
from xgboost import XGBClassifier
clf_xgb = XGBClassifier(n_estimators=100, random_state=42)
clf_xgb.fit(X_train, y_train)
y_pred = clf_xgb.predict(X_test)
accuracy_xgb = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy_xgb:.2f}')
print('\nClassification Report:\n', classification_report(y_test, y_pred))
```


Accuracy: 0.99

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	333
1	0.95	0.95	0.95	64
2	1.00	0.97	0.98	29
accuracy			0.99	426
macro avg	0.98	0.97	0.98	426
weighted avg	0.99	0.99	0.99	426

```
cf= confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix for XGB')
```

Text(0.5, 1.0, 'Confusion Matrix for XGB')



```
from sklearn.svm import SVC
clf_svm = SVC(kernel='linear', C=1.0, random_state=42)
clf_svm.fit(X_train, y_train)
y_pred_svm = clf_svm.predict(X_test)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print(f'SVM Accuracy: {accuracy_svm:.2f}')
print('\nSVM Classification Report:\n', classification_report(y_test, y_pred_svm))
```

SVM Accuracy: 0.98


SVM Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	333
1	0.95	0.92	0.94	64
2	1.00	0.97	0.98	29
accuracy			0.98	426

macro avg	0.98	0.96	0.97	426
weighted avg	0.98	0.98	0.98	426

```
cf= confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix for SVM')
```

Text(0.5, 1.0, 'Confusion Matrix for SVM')



```
from sklearn.neighbors import KNeighborsClassifier
clf_knn = KNeighborsClassifier(n_neighbors=5)
clf_knn.fit(X_train, y_train)
y_pred_knn = clf_knn.predict(X_test)
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print(f'KNN Accuracy: {accuracy_knn:.2f}')
print('\nKNN Classification Report:\n', classification_report(y_test, y_pred_knn))
```


KNN Accuracy: 0.92

KNN Classification Report:				
	precision	recall	f1-score	support
0	0.94	0.97	0.95	333
1	0.83	0.67	0.74	64
2	0.96	0.93	0.95	29
accuracy			0.92	426
macro avg	0.91	0.86	0.88	426
weighted avg	0.92	0.92	0.92	426

```
cf= confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix for KNN')
```

```
Text(0.5, 1.0, 'Confusion Matrix for KNN')
```

Confusion Matrix for KNN



```
from lightgbm import LGBMClassifier
clf_lgbm = LGBMClassifier(n_estimators=100, random_state=42)
clf_lgbm.fit(X_train, y_train)
y_pred_lgbm = clf_lgbm.predict(X_test)
accuracy_lgbm = accuracy_score(y_test, y_pred_lgbm)
print(f'LGBM Accuracy: {accuracy_lgbm:.2f}')
print('\nLGBM Classification Report:\n', classification_report(y_test, y_pred_lgbm))
```

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000387 seconds. You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[illegible]

LGBM Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	333
1	0.95	0.94	0.94	64
2	1.00	0.97	0.98	29
accuracy			0.98	426
macro avg	0.98	0.97	0.97	426
weighted avg	0.98	0.98	0.98	426

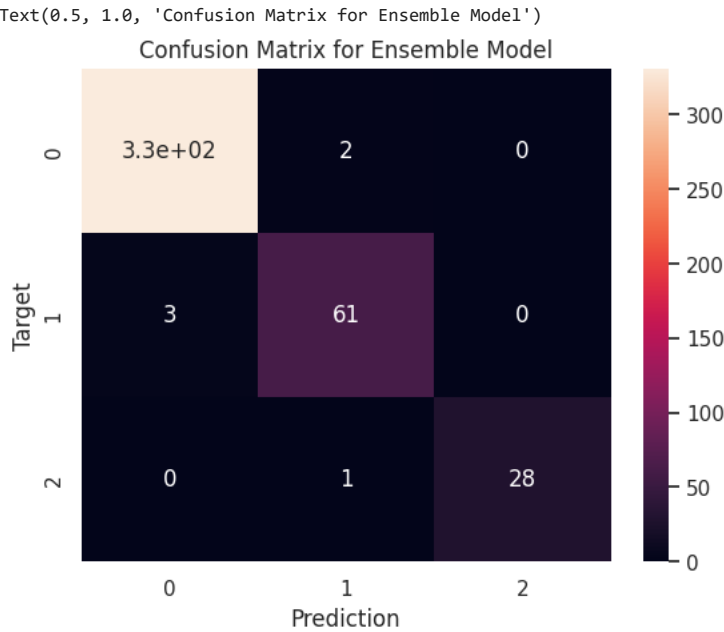
```
cf= confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix for LGBM')
```


Confusion Matrix for LGBM

	Actual 0	Actual 1
Predicted 0	3.3e+02	2
Predicted 1	0	1

[illegible]

```
cf= confusion_matrix(y_test, y_pred)
plt.figure()
sns.heatmap(cf, annot=True)
plt.xlabel('Prediction')
plt.ylabel('Target')
plt.title('Confusion Matrix for Ensemble Model')
```



```
accuracies = [  
    ('Random Forest', accuracy_rf),  
    ('XGBoost', accuracy_xgb),  
    ('SVM', accuracy_svm),  
    ('KNN', accuracy_knn),  
    ('LGBM', accuracy_lgbm),  
    ('Ensemble', accuracy_ensemble)  
]  
  
accuracies_df = pd.DataFrame(accuracies, columns=['Classifier', 'Accuracy'])  
plt.figure(figsize=(10, 6))  
sns.set(style="whitegrid")  
sns.barplot(x='Classifier', y='Accuracy', data=accuracies_df, palette="viridis")  
plt.title('Classifier Accuracies Comparison')  
plt.yticks([i/10 for i in range(11)]) # Set y-axis ticks in intervals of 0.1  
plt.ylim(0, 1)  
plt.show()
```

